

C++11 Move

l-value / r-value c++03

```
int fct(int const & a) {  
    return a;  
}  
  
int main() {  
    // the right way: lvalues left of = / rvalues right of =  
    int l = 10;  
    int const lc = fct(l);  
  
    // the wrong way: rvalues left of =  
    fct(l) = l;  
    10 = lc;  
  
    // lvalues may also be on the right side  
  
    return 0;  
}
```

the valueness is **independent** of the type
(i.e. there are `int` lvalues and `int` rvalues)

l-value

- “can be on left of =”
- Has an identity
- Examples:
 - `int lval;`
 - `int const cval;`
 - `++lval;`
 - `std::cout;`
 - `std::cout << lval;`
 - `a=b;`
 - `a+=b;`
 - `fct(); // int & fct()`

r-value

- “can only be on right of =”
- Does not have an identity
- Examples:
 - `1+2;`
 - `lval++;`
 - `fct(); // int fct()`
 - `a == b;`
 - `&lval;`

r-value references

```
int main() {  
  
    int l = 10;  
  
    int & lref = l;  
  
    int & lref = 10; // fails!  
    int const & clref = 10; // ok  
  
    // the new c++11 rvalue reference  
    int && rref = 10; // works  
  
    // so what have we gained?  
    // Not much without move-constructors / move-assignment  
  
    return 0;  
}
```

„Factory“ Function Problem

```
std::vector<int> all_int_prime_nr() { std::vector<int> res; ... return res; };  
void all_int_prime_nr_ref(std::vector<int> & pnr);
```

```
int main() {  
    // looks like it needs one ctor, one copy-ctor and one dtor  
    // don't we have at some point two (huge) copies?  
    // actually, this is no problem due to „copy elision“, only one ctor  
    std::vector<int> pnr = all_int_prime_nr();  
  
    // only one ctor, perhaps more efficient  
    std::vector<int> pnr_efficient;  
    all_int_prime_nr_ref(pnr_efficient);  
  
    // but as soon as there is no construction anymore the problem arises  
    std::vector<int> pnr_later;  
    pnr_later = all_int_prime_nr();  
  
    return 0;  
}
```

std::vector (c++03)

```
template<typename T, typename Alloc = allocator<T> >
class vector {

    public:
        using size_type = uint32_t;
        using value_type = T;

        vector();
        ~vector();
        vector(vector const & rhs);
        vector & operator=(vector const & rhs);

        void push_back(value_type const & val);
        void reserve(size_type const & new_cap);
        void clear();

        size_type const & size() const { return size_; }
        size_type const & capacity() const { return capacity_; };

    private:
        void check_capacity(); // small helper
        value_type * data_;
        size_type size_;
        size_type capacity_;
};
```

std::vector structors

```
// constructor
vector(): data_(nullptr)
        , size_(0)
        , capacity_(0)
{}


// copy-constructor
vector(vector const & rhs): data_(Alloc::allocate(rhs.size()))
        , size_(0)
        , capacity_(rhs.size()) {

    for(size_type i = 0; i < rhs.size(); ++i)
        Alloc::construct(data_ + size_++, rhs[i]);
}

// destructor
~vector() {
    clear();
    Alloc::deallocate(data_, capacity());
}
```

push_back (c++03)

```
void push_back(value_type const & val) {  
    check_capacity();  
    Alloc::construct(data_ + size_++, val);  
}  
void check_capacity() {  
    if(size() == capacity()) {  
        if(capacity() == 0)  
            reserve(1);  
        else  
            reserve(2*capacity());  
    }  
}  
void reserve(size_type const & new_cap) {  
    value_type * new_data = Alloc::allocate(new_cap);  
  
    for(size_type i = 0; i < size(); ++i)  
        Alloc::construct(new_data + i, data_[i]);  
  
    for(size_type i = 0; i < size(); ++i)  
        Alloc::destruct(data_ + i);  
  
    Alloc::deallocate(data_, capacity());  
    data_ = new_data;  
    capacity_ = new_cap;  
}
```



two loops for the strong
exception guarantee

allocator (c++03)

```
template<typename T>
struct allocator {

    static T * allocate(std::size_t const & n) {
        // this form of new returns n * sizeof(T) bytes memory
        return static_cast<T*>(::operator new(n * sizeof(T)));
    }
    static void deallocate(T * p, std::size_t const & n) {
        ::operator delete(p);
    }

    static void construct(T * p, T const & t) {
        // this is called a „placement new“. It doesn't allocate any memory
        new (p) T(t); // copy-ctor call
    }
    static void destruct(T * p) { p->~T(); }
};
```

what happens in memory

```
int main() {  
    vector<myint> a;
```

```
    return 0;
```

```
}
```

a.data_ = nullptr								a.size_ = 0				a.cap_ = 0			

```
struct myint {  
    // ctor  
    myint(int const & a): x(a) {}  
  
    // copy  
    myint(myint const &);  
    myint & operator=(myint const &);  
  
    int8_t x;  
};
```


what happens in memory

```
int main() {  
    vector<myint> a;  
  
    a.push_back(0);  
}
```

```
    return 0;
}
```

[illegible]


what happens in memory

```
int main() {  
    vector<myint> a;
```

```
a.push_back(0);  
a.push_back(1);
```

```
return 0;
```

}




what happens in memory

```
int main() {  
    vector<myint> a;
```

```
    a.push_back(0);  
    a.push_back(1);  
    a.push_back(2);  
    a.push_back(3);
```

```
    return 0;
```

```
}
```



what happens in memory

```
int main() {  
    vector<myint> a;
```

```
a.push_back(0);
a.push_back(1);
a.push_back(2);
a.push_back(3);
a.push_back(4);
```

```
return 0;
```

}

[illegible]


what happens in memory

```
int main() {  
    vector<myint> a;
```

```
    a.push_back(0);  
    a.push_back(1);  
    a.push_back(2);  
    a.push_back(3);  
    a.push_back(4);
```

```
    return 0;
```

```
}
```




what happens in memory

```
int main() {  
    vector<myint> a;
```

```
    a.push_back(0);  
    a.push_back(1);  
    a.push_back(2);  
    a.push_back(3);  
    a.push_back(4);
```

```
    return 0;
```

```
}
```



what happens in memory

```
int main() {
    vector<myint> a;
```

```
a.push_back(0);
a.push_back(1);
a.push_back(2);
a.push_back(3);
a.push_back(4);
```



```
return 0;
```

}

this needed (regarding `myint`):

- 15 byte allocation
- 7 bytes deallocation
- 5 ctor calls
- 7+5 copy-ctor calls
- 7+5 dtor calls

why do we have these ops?


- 5 copy-ctor calls
- 5 dtor calls

[illegible]

additional temporary copy

```
int main() {  
    vector<myint> a;  
  
    a.push_back(0);  
}
```

```
// if an int is passed, val is a temporary object
void push_back(value_type const & val) {
    check_capacity();
```



```
return 0;
```

}


[illegible]

additional temporary copy

```
int main() {  
    vector<myint> a;  
  
    a.push_back(0);
```

```
// if an int is passed, val is a temporary object  
void push_back(value_type const & val) {  
    check_capacity();  
    Alloc::construct(data_ + size_++, val);  
}
```

```
    return 0;  
}
```



we can remove these additional copy-ctor calls next week

std::vector structors

```
// constructor
```

```
vector(): data_(nullptr)
        , size_(0)
        , capacity_(0)
{}

```

```
// copy-constructor
```

```
vector(vector const & rhs): data_(Alloc::allocate(rhs.size()))
        , size_(0)
        , capacity_(rhs.size()) {

    for(size_type i = 0; i < rhs.size(); ++i)
        Alloc::construct(data_ + size_++, rhs[i]);
}

```

```
// destructor
```

```
~vector() {
    clear();
    Alloc::deallocate(data_, capacity());
}


```

copy-ctor

```
int main() {  
    ...  
    // copy-ctor is called  
    vector<myint> b(a);  
    return 0;  
}
```

this needed (regarding myint):

- a.size() byte allocation
- a.size() copy-ctor calls



why move?

- copying from A to B and later delete A (like in `reserve`) seems inefficient
- we want to reuse A's resources to construct B
- some entities are non-copyable (i.e. threads) but they are usually moveable

std::vector move ctor

```
// constructor
vector(): data_(nullptr)
        , size_(0)
        , capacity_(0)
{}

// move-constructor (note: rhs is not const)
vector(vector && rhs):
        data_(rhs.data_)
        , size_(rhs.size_)
        , capacity_(rhs.capacity_)
{
    // rhs's invariants must not be broken!
    // without the following lines we only have a shallow copy...
    rhs.data_ = nullptr;
    rhs.size_ = 0;
    rhs.capacity_ = 0;
}
```

semantic convention: the moved-from object must be left in a valid state (i.e. not break it's invariants)

move-ctor

```
int main() {  
    ...  
    // move-ctor is called  
    vector<myint> b(std::move(a));  
    return 0;  
}
```

`std::move` does not move anything! it casts the object to an r-value reference, so that the overload mechanism prefers the move-ctor (better match)

a.data_ = nullptr						a.size_ = 0				a.cap_ = 0			
b.data_						a.size_ = 5				a.cap_ = 8			
	0	1	2	3	4								

this needed (regarding myint):

- 0 byte allocation
- 0 copy-ctor calls

std::move implementation

```
template<typename T>
constexpr std::remove_reference_t<T> && move(T && t) noexcept {
    return static_cast<std::remove_reference_t<T> &&>(t);
}
```

std::move does not move anything! it casts the object to an r-value reference, so that the overload mechanism prefers the move-ctor (better match)

std::vector move assign

```
// constructor
vector(): data_(nullptr)
        , size_(0)
        , capacity_(0)
{}

// move-assignment (note: rhs is not const)
vector & operator=(vector && rhs) {
    using std::swap;
    swap(data_, rhs.data_);
    swap(size_, rhs.size_);
    swap(capacity_, rhs.capacity_);
    rhs.clear(); // semantically optional, the std::vector clears the rhs

    return (*this);
}
```

move-assign

```
int main() {  
    ...  
    // move-assign is called  
    vector<myint> b;  
    b = std::move(a);  
    return 0;  
}
```

a.data_ = nullptr						a.size_ = 0				a.cap_ = 0			
b.data_						a.size_ = 5				a.cap_ = 8			
						this needed (regarding myint): • b.size() dtor calls							
	0	1	2	3	4								

conclusion

move operation are almost always (much) faster than their corresponding copy operation!

Especially for classes that contain dynamically allocated data (e.g. vector / pimpl)

improve reserve with move

```
void push_back(value_type const & val) {
    check_capacity();
    Alloc::construct(data_ + size_++, val);
}

void check_capacity() {
    if(size() == capacity()) {
        if(capacity() == 0)
            reserve(1);
        else
            reserve(2*capacity());
    }
}

void reserve(size_type const & new_cap) {
    value_type * new_data = Alloc::allocate(new_cap);

    for(size_type i = 0; i < size(); ++i)
        Alloc::construct(new_data + i, std::move(data_[i]));

    for(size_type i = 0; i < size(); ++i)
        Alloc::destruct(data_ + i);

    Alloc::deallocate(data_, capacity());
    data_ = new_data;
    capacity_ = new_cap;
}
```

is this enough to enable
the move-ctor instead of
the copy-ctor?

improve allocator with move

```
template<typename T>
struct allocator {

    static T * allocate(std::size_t const & n);
    static void deallocate(T * p, std::size_t const & n);

    static void construct(T * p, T const & t) {
        // this is called a „placement new“. It doesn't allocate any memory
        new (p) T(t);
    }
    static void construct(T * p, T && t) {
        new (p) T(std::move(t));
    }

    static void destruct(T * p) { p->~T(); }
};
```

is this enough now to use
the move-ctor of myint
instead of the copy-ctor?

myint needs move support

```
struct myint {  
    // ctor  
    myint(int const & a): x(a) {}  
  
    // copy (disables move)  
    myint(myint const &);  
    myint & operator=(myint const &);
```

```
    // move  
    myint(myint &&);  
    myint & operator=(myint &&);
```

```
    int8_t x;  
};
```

		compiler generates (if possible)					
		default ctor	copy ctor	copy assign	move ctor	move assign	dtor
w e d e c l a r e	<i>any ctor</i>	no	yes	yes	yes	yes	yes
	<i>copy ctor</i>	yes		yes*	no	no	yes
	<i>copy assign</i>	yes	yes*		no	no	yes
	<i>move ctor</i>	yes	no	no		no	yes
	<i>move assign</i>	yes	no	no	no		yes
	<i>dtor</i>	yes	yes*	yes*	no	no	

yes* should be no (backwards compatible)


what happens in memory

```
int main() {
    vector<myint> a;
```

```
a.push_back(0);
a.push_back(1);
a.push_back(2);
a.push_back(3);
a.push_back(4);
```

```
return 0;
```

}



what happens in memory

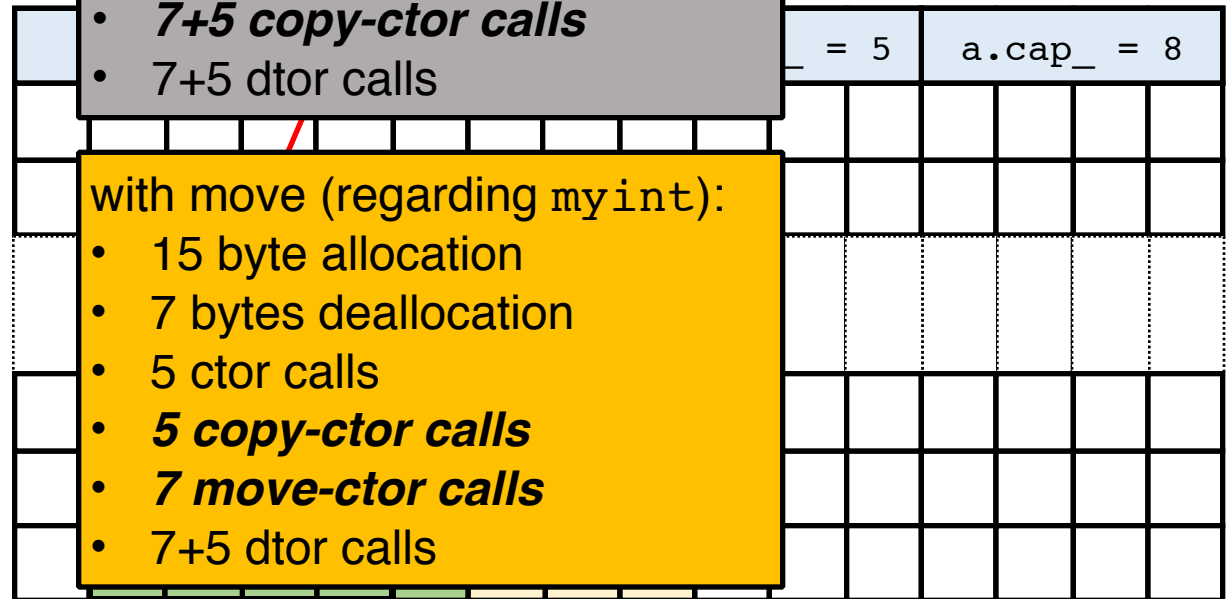
```
int main() {  
    vector<myint> a;  
  
    a.push_back(0);  
    a.push_back(1);  
    a.push_back(2);  
    a.push_back(3);  
    a.push_back(4);  
  
    return 0;  
}
```

before with copy only:

- 15 byte allocation
- 7 bytes deallocation
- 5 ctor calls
- **7+5 copy-ctor calls**
- 7+5 dtor calls

with move (regarding myint):

- 15 byte allocation
- 7 bytes deallocation
- 5 ctor calls
- **5 copy-ctor calls**
- **7 move-ctor calls**
- 7+5 dtor calls



exception safety in reserve

```
void push_back(value_type const & val) {
    check_capacity();
    Alloc::construct(data_ + size_++, val);
}

void check_capacity() {
    if(size() == capacity()) {
        if(capacity() == 0)
            reserve(1);
        else
            reserve(2*capacity());
    }
}

void reserve(size_type const & new_cap) {
    value_type * new_data = Alloc::allocate(new_cap);

    for(size_type i = 0; i < size(); ++i)
        Alloc::construct(new_data + i, std::move(data_[i]));

    for(size_type i = 0; i < size(); ++i)
        Alloc::destruct(data_ + i);

    Alloc::deallocate(data_, capacity());
    data_ = new_data;
    capacity_ = new_cap;
}
```

why is the strong exception guarantee broken?

exception safety in reserve

```
void push_back(value_type const & val) {
    check_capacity();
    Alloc::construct(data_ + size_++, val);
}

void check_capacity() {
    if(size() == capacity()) {
        if(capacity() == 0)
            reserve(1);
        else
            reserve(2*capacity());
    }
}

void reserve(size_type const & new_cap) {
    value_type * new_data = Alloc::allocate(new_cap);

    for(size_type i = 0; i < size(); ++i)
        Alloc::construct(new_data + i, std::move_if_noexcept(data_[i]));

    for(size_type i = 0; i < size(); ++i)
        Alloc::destruct(data_ + i);

    Alloc::deallocate(data_, capacity_);
    data_ = new_data;
    capacity_ = new_cap;
}
```

“move if you can - copy if you must” strategy

noexcept

```
// does not throw
// we (!) promis the compiler with noexcept that this fct does not throw
int f1(int const & a) noexcept {
    return a;
}
// could still be declared noexcept, but the compiled code
// could erase your hard disk if something is thrown ;)
int f2(int const & a) {
    throw std::runtime_error("never call this");
    return a;
}
// exception neutral function
// doesn't throw itself but calls something that may
int g1(int const & a) noexcept(noexcept(f1(a))) {
    return f1(a);
}
int g2(int const & a) noexcept(noexcept(f2(a))) {
    return f2(a);
}
int main() {
    std::cout << noexcept(f1(1)) << std::endl; // 1
    std::cout << noexcept(g1(1)) << std::endl; // 1
    std::cout << noexcept(f2(1)) << std::endl; // 0
    std::cout << noexcept(g2(1)) << std::endl; // 0
}
```

noexcept

```
// does not throw
// we (!) promise the compiler with noexcept that this fct does not throw
int f1(int const & a) noexcept {
    return a;
}
// could still be declared noexcept, but the compiled code
// could erase your hard disk if something is thrown ;)
int f2(int const & a) {
    throw std::runtime_error("never call this");
    return a;
}
// exception neutral function
// doesn't throw itself but calls s
int g1(int const & a) noexcept(noexcept(f1(a))) {
    return f1(a);
}
int g2(int const & a) noexcept(noexcept(f2(a))) {
    return f2(a);
}
int main() {
    std::cout << noexcept(f1(1)) << std::endl; // 1
    std::cout << noexcept(g1(1)) << std::endl; // 1
    std::cout << noexcept(f2(1)) << std::endl; // 0
    std::cout << noexcept(g2(1)) << std::endl; // 0
}
```

noexcept allows for:

- many compiler optimizations
- better exception guarantees

noexcept

```
struct myint {  
    // ctor  
    myint(int const & a): x(a) {}  
  
    // copy (disables move)  
    myint(myint const &) noexcept;  
    myint & operator=(myint const &) noexcept;  
  
    // move  
    myint(myint const &) noexcept;  
    myint & operator=(myint const &) noexcept;  
  
    int8_t x;  
};
```

use noexcept wherever possible. It may allows libraries to run faster/safer versions of their implementation.