

Boost.Serialization

– a short introduction –

Developed by Robert Ramey
<http://www.rrsd.com/>



Serialization

- ◆ the reversible deconstruction of an arbitrary set of C++ data structures to a sequence of bytes
- ◆ Can be used for
 - ◆ Storing to and loading from a file
 - ◆ Remote parameter passing
 - ◆ Message passing
 - ◆ ...

Archive

◆ a specific rendering of this stream of bytes.

◆ Examples

◆ Binary file

◆ Text file

◆ XML file

◆ Buffer in memory

◆ MPI message

◆ ...

Requirements

- ◆ Non-intrusive
- ◆ Versioning for each class
- ◆ Deep pointer save and restore
- ◆ Proper restoration of pointers to shared data
- ◆ Enable data portability
- ◆ Orthogonal specification of class serialization and archive format

A first example

```
class Client
{
public:
    Client()
    {}

    Client(int c, std::string r)
        : cif_number(c),
          rating(r)
    {}

private:
    int cif_number;
    std::string rating;
};
```

```
std::ofstream ofs("filename");
boost::archive::text_oarchive oa(ofs);

const Client c(123521, "AAA");
oa << c;
ofs.close();

_____

std::ifstream ifs("filename");
boost::archive::text_iarchive ia(ifs);

Client c;
ia >> c;
ifs.close();
```

Boost.Serialization overview

◆ Serialization of

- ◆ Primitive types
- ◆ Arrays
- ◆ Pointers
- ◆ Classes
- ◆ Derived Classes through base class pointers
- ◆ ...

◆ Archive types

- ◆ Text archive
- ◆ Wide character text archive
- ◆ Binary archive
- ◆ XML archive
- ◆ Wide character XML archive

◆ Additional archive types are not hard to implement

Boost.Serialization archives

- ◆ A (nearly) portable text archives
 - ◆ `boost::archive::text_oarchive(ostream &s)`
 - ◆ `boost::archive::text_iarchive(istream &s)`
 - ◆ `boost::archive::text_woarchive(wostream &s)`
`boost::archive::text_wiarchive(wistream &s)`
- ◆ Non-portable native binary archives
 - ◆ `boost::archive::binary_oarchive(ostream &s)`
 - ◆ `boost::archive::binary_iarchive(istream &s)`
- ◆ XML archives
 - ◆ `boost::archive::xml_oarchive(ostream &s)`
 - ◆ `boost::archive::xml_iarchive(istream &s)`
 - ◆ `boost::archive::xml_woarchive(wostream &s)`
 - ◆ `boost::archive::xml_wiarchive(wistream &s)`

Intrusive Implementation

```
class Client
{
public:
    Client();
    Client(int c, std::string r);

private:
    int cif_number;
    std::string rating;

friend class boost::serialization::access;

template<class Archive>
void serialize(Archive & ar, const unsigned int version)
{
    ar & cif_number & rating;
}
};
```


Less Intrusive Implementation

```
class Client
{
public:
    Client();
    Client(int c, std::string r);

public:
    int cif_number;
    std::string rating;
};
```

```
template<class Archive>
void serialize(
    Archive & ar,
    Client & g,
    const unsigned int version
)
{
    ar & g.cif_number;
    ar & g.rating;
}
```

Derived classes

```
class PrivatePerson : public Client
{
public:
    PrivatePerson();
    PrivatePerson(int c, std::string r, std::string n);

private:
    std::string name;

    friend class boost::serialization::access;

    template<class Archive>
    void serialize(Archive & ar, const unsigned int version) {
        ar & boost::serialization::base_object< Client >(*this);
        ar & name;
    }
};
```

Serializing Pointers Works

```
boost::archive::text_oarchive oa(std::ofstream("filename"));
```

```
const Client* cp1 = new Client (123521,"OK");
```

```
const Client* cp2 = cp1;
```

```
oa << cp1;          // This serializes the object and pointer
```

```
oa << cp2;          // This serializes only the pointer
```

◆ When serializing objects

- ◆ they are assigned an integer identifier
- ◆ their address gets registered

◆ When serializing a pointer

- ◆ The pointee is serialized if it has not been serialized yet
- ◆ The identifier of the pointee is written to the archive

Deserializing Pointers Works

```
boost::archive::text_iarchive ia(std::ifstream("filename"));
```

```
const Client* cp1 , cp2;
```

```
ia >> cp1;          // creates the object and deserializes the object and pointer  
ia >> cp2;          // deserializes the pointer, pointing to the same object
```

- ◆ When deserializing objects
 - ◆ they are assigned an integer identifier
 - ◆ their address gets registered
- ◆ When deserializing a pointer
 - ◆ The identifier of the pointee is read from the archive
 - ◆ The pointee is deserialized if it has not been deserialized yet
 - ◆ The pointer is set

Serializing Arrays

```
class ClientArray
{
public:
    ClientArray();

private:
    Client * clients[1000];

    friend class boost::serialization::access;

    template<class Archive>
    void serialize(Archive & ar, const unsigned int version)
    {
        ar & clients;
    }
};
```

Serializing STL Collections

```
class ClientArray
{
public:
    ClientArray();

private:
    std::vector<Client*> clients;

    friend class boost::serialization::access;

    template<class Archive>
    void serialize(Archive & ar, const unsigned int version)
    {
        ar & clients;
    }
};
```

Calculated members

```
class FixedRate
{
public:
    FixedRate(double a, double r);

private:
    double amount, rate, interest;

    friend class boost::serialization::access;
    template<class Archive>
    void serialize(Archive & ar, const unsigned int)
    {
        ar & amount & rate & interest; // wastes space in archive
    }
};
```

```
FixedRate::FixedRate(double a, double r)
    : amount(a),
      rate(r),
      interest(a*r)
{ }
```

Calculated members

```
class FixedRate
{
public:
    FixedRate(double a=0., double r=0.);

private:
    double amount, rate, interest;

    friend class boost::serialization::access;
    template<class Archive>
    void serialize(Archive & ar, const unsigned int)
    {
        ar & amount & rate; // fails to recalculate interest on loading
    }
};
```


Splitting serialize

```
class FixedRate
{
public:
    ...
private:
    ...
    template<class Archive>
    void save(Archive & ar, const unsigned int) const
    {
        ar & amount & rate;           // serialize amount and rate
    }

    template<class Archive>
    void load(Archive & ar, const unsigned int)
    {
        ar & amount & rate;           // deserialize amount and rate
        interest = amount*rate;      // calculate interest
    }

    BOOST_SERIALIZATION_SPLIT_MEMBER()
};
```

Pointers to Objects of Derived Classes

```
class base
```

```
{
```

```
    ...
```

```
};
```

```
class derived_one : public base
```

```
{
```

```
    ...
```

```
};
```

```
class derived_two : public base
```

```
{
```

```
    ...
```

```
};
```

```
int main()
```

```
{
```

```
    ...
```

```
    base *b = new derived_one();
```

```
    oa << b; // which derived class?
```

```
}
```

```
int main()
```

```
{
```

```
    ...
```

```
    base *b;
```

```
    ia >> b; // which derived class?
```

```
}
```

Registering derived types

◆ Need to register all derived types

```
#include <boost/serialization/export.hpp>
```

```
BOOST_CLASS_EXPORT(derived_one)
```

```
BOOST_CLASS_EXPORT(derived_two)
```

◆ Note

- ◆ Serialization of derived classes **must** call base class serialization even if that is empty

XML Serialization

```
class Client
{
public:
    Client(
    Client(int c, std::string r);


private:
    int number;
    std::string rating;

    friend class boost::serialization::access;
    template<class Archive>
    void serialize(Archive & ar, const unsigned version)
    {
        ar & BOOST_SERIALIZATION_NVP(number);
        ar & BOOST_SERIALIZATION_NVP(rating);
    }
};
```

→

```
<item class_id="0" ...>
    <number>132542</number>
    <rating>bankrupt</rating>
</item>
```

XML Serialization

<pre>class PrivatePerson : public Client { public: PrivatePerson(); PrivatePerson(int, std::string, std::string); private: std::string name; friend class boost::serialization::access; template<class Archive> void serialize(Archive & ar, const unsigned version) { ar & BOOST_SERIALIZATION_BASE_OBJECT_NVP< Client >(*this); ar & BOOST_SERIALIZATION_NVP(name); } };</pre>		<pre><item class id="0" ...> <Client class_id="1" ...> <number>132542</number> <rating>bankrupt</rating> </Client> <name>Dave Abrahams</name> </item></pre>
--	--	---

Serialization traits: Implementation level

- ◆ Specified by specializing the **level** template, or macro call:

BOOST_CLASS_IMPLEMENTATION(my_class, *value*)

- ◆ Possible values

- ◆ boost::serialization::**not_serializable**

-> do not serialize

- ◆ boost::serialization::**primitive_type**

-> archive knows how to serialize the type

- ◆ boost::serialization::**object_serializable**

-> call the object's serialize function

- ◆ boost::serialization::**object_class_info**

-> store class info (version, name) and call the object's serialize function

Object tracking

- ◆ Determines whether object addresses should be registered to allow serialization of pointers
- ◆ Specified by specializing the **tracking** template, or macro call:

BOOST_CLASS_TRACKING(my_class, *value*)

- ◆ Possible values
 - ◆ `boost::serialization::track_never`
-> never track the object's address
 - ◆ `boost::serialization::track_selectively`
-> track only if the type was explicitly registered
 - ◆ `boost::serialization::track_always`
-> track always

New Versions of Classes

```
class Client
{
    ...
private:
    int cif_number;
    std::string rating;

    friend class boost::serialization::access;

    template<class Archive>
    void serialize(Archive & ar, const unsigned int version)
    {
        ar & number & rating;
    }
};
```


New Versions of Classes

```
class Client
{
    ...
private:
    int cif_number;
    std::string rating;
    bool simulated_client;
    friend class boost::serialization::access;

    template<class Archive>
    void serialize(Archive & ar, const unsigned int version)
    {
        ar & number & rating & simulated_client;
    }
};
```

- We can no longer read the old files!

Class Versioning

```
class Client
{
    ...
private:
    int cif_number;
    std::string rating;
    bool simulated_client;
    friend class boost::serialization::access;

    template<class Archive>
    void serialize(Archive & ar, const unsigned int version)
    {
        ar & number & rating;
        if (version > 0)
            ar & simulated_client
    }
};
```

BOOST_CLASS_VERSION(Client, 1)

Serialization traits: version

- ◆ Specified by specializing the **version** template

```
namespace boost { namespace serialization {
```

```
    struct version<my_class>
    {
        static const unsigned int value = 2;
    };

}}
```

- ◆ Convenience macro

```
BOOST_CLASS_VERSION(my_class, 2)
```

Classes without default constructor

```
Class Counterparty {  
    public:  
        Counterparty(int id)  
            : identifier(id)  
        {}  
    private:  
        int identifier;  
}
```

```
Counterparty *c;  
ia >> c; // fails!  
  
// Boost.serialization attempts:  
  
c = new Counterparty(); // fails!  
ia >> *c;  
  
// Boost.serialization should do:  
  
int id;  
ia >> id;  
c = new Counterparty(id);
```

Overload construction of object

```
namespace boost { namespace serialization {  
  
template<class Archive>  
inline void load_construct_data  
    (  
        Archive & ar,  
        Counterparty * c,  
        const unsigned version  
    )  
{  
    int m;  
    ar >> m; // retrieve data from archive required to construct new instance  
    ::new(c) Counterparty(m); // invoke inplace constructor to initialize  
}  
  
}}
```

Serializing reference members

- ◆ Similar problem as before since reference members usually require non-default constructor
- ◆ Solution: serialize references as pointers and overload **save_construct_data** and **load_construct_data**
- ◆ Look at the documentation for details

Advanced topics

◆ See the documentation for information on

- ◆ Creating your own archives
- ◆ Serialization exceptions
- ◆ Polymorphic archives
- ◆ Portable archives
- ◆ ...